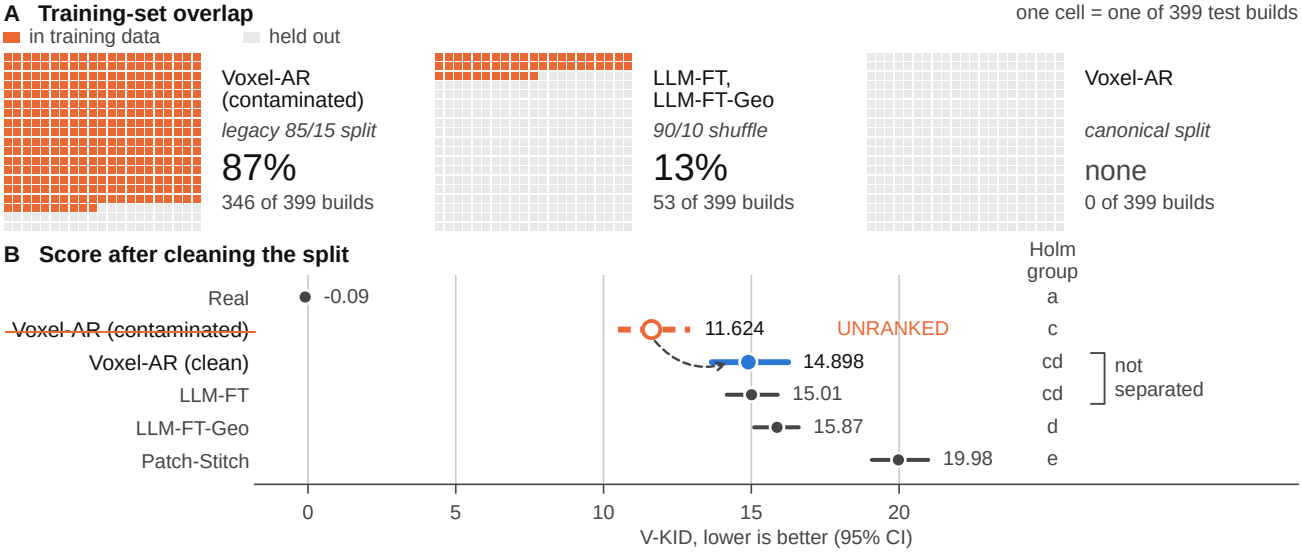
A Training-set overlap
one cell = one of 399 test builds
in training data
held out
Voxel-AR (contaminated)
legacy 85/15 split
87%
346 of 399 builds
LLM-FT, LLM-FT-Geo
90/10 shuffle
13%
53 of 399 builds
Voxel-AR
canonical split
none
0 of 399 builds
B Score after cleaning the split
Holm group
Real
-0.09
a
Voxel-AR (contaminated)
11.624
UNRANKED
c
Voxel-AR (clean)
14.898
cd
not separated
LLM-FT
15.01
cd
LLM-FT-Geo
15.87
d
Patch-Stitch
19.98
e
0
5
10
15
20
V-KID, lower is better (95% CI)