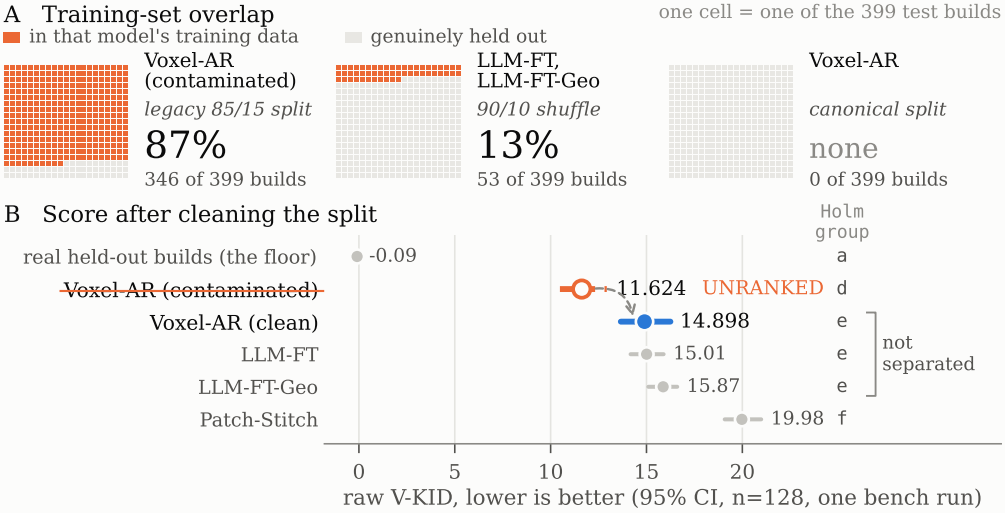
A Training-set overlap
in that model's training data
genuinely held out
one cell = one of the 399 test builds
Voxel-AR (contaminated)
legacy 85/15 split
87%
346 of 399 builds
LLM-FT, LLM-FT-Geo
90/10 shuffle
13%
53 of 399 builds
Voxel-AR
canonical split
none
0 of 399 builds
B Score after cleaning the split
Holm group
real held-out builds (the floor)
-0.09
a
Voxel-AR (contaminated)
11.624
UNRANKED
d
Voxel-AR (clean)
14.898
e
LLM-FT
15.01
e
not separated
LLM-FT-Geo
15.87
e
Patch-Stitch
19.98
f
0
5
10
15
20
raw V-KID, lower is better (95% CI, n=128, one bench run)